

# MAGI Algorithm Description

Kamil Hlubek

June 16, 2025

## Notation

- $\mathcal{I} = \{t_1, \dots, t_n\}$ : Discretization time grid.
- $D$ : Number of dimensions (state variables).
- $n$ : Number of discretization time points.
- $k$ : Number of ODE parameters.
- $N_{obs}$ : Number of distinct observation time points (may be less than  $n$ ).
- $Y \in \mathbb{R}^{n \times D}$ : Observation matrix on grid  $\mathcal{I}$ , with NaNs for unobserved points/times.  $Y_{j,d}$  is observation for dimension  $d$  at time  $t_j$ . Let  $Y_d = Y_{:,d}$  be the vector for dimension  $d$ .
- $X \in \mathbb{R}^{n \times D}$ : Latent state matrix on grid  $\mathcal{I}$ .  $X_{j,d} = x_d(t_j)$ . Let  $X_d = X_{:,d}$  be the vector for dimension  $d$ .
- $\theta \in \Theta \subset \mathbb{R}^k$ : Vector of ODE parameters.  $\Theta$  defines the parameter bounds  $[\theta_{lower}, \theta_{upper}]$ .
- $\sigma \in \mathbb{R}_{>0}^D$ : Vector of observation noise standard deviations for each dimension.  $\sigma^2 = (\sigma_1^2, \dots, \sigma_D^2)$ .
- $\phi_d = (\phi_{d,1}, \phi_{d,2})$ : GP hyperparameters (variance  $\phi_{d,1}$ , lengthscale  $\phi_{d,2}$ ) for dimension  $d$ .  $\phi \in \mathbb{R}_{>0}^{2 \times D}$ .
- $\mathcal{K}_{\phi_d}(t, t')$ : GP covariance kernel function for dimension  $d$ .
- $f(x, \theta, t) : \mathbb{R}^D \times \mathbb{R}^k \times \mathbb{R} \rightarrow \mathbb{R}^D$ : ODE function,  $dx/dt = f(x, \theta, t)$ .  $f_d$  is the  $d$ -th component. Let  $F_d \in \mathbb{R}^n$  be the vector where  $[F_d]_j = f_d(X_{j,:}, \theta, t_j)$ .
- $J_x(x, \theta, t) = \nabla_x f(x, \theta, t) \in \mathbb{R}^{D \times D}$ : Jacobian of  $f$  w.r.t. state  $x$ .
- $J_\theta(x, \theta, t) = \nabla_\theta f(x, \theta, t) \in \mathbb{R}^{D \times k}$ : Jacobian of  $f$  w.r.t. parameters  $\theta$ .
- $C_{\phi_d} \in \mathbb{R}^{n \times n}$ : GP prior covariance matrix for dim  $d$ ,  $[C_{\phi_d}]_{jj'} = \mathcal{K}_{\phi_d}(t_j, t_{j'})$ .
- $C'_{\phi_d} \in \mathbb{R}^{n \times n}$ : First time derivative matrix,  $[C'_{\phi_d}]_{jj'} = \frac{\partial}{\partial t} \mathcal{K}_{\phi_d}(t, t')|_{t=t_j, t'=t_{j'}}$ .
- $C''_{\phi_d} \in \mathbb{R}^{n \times n}$ : Second time derivative matrix,  $[C''_{\phi_d}]_{jj'} = \frac{\partial^2}{\partial t \partial t'} \mathcal{K}_{\phi_d}(t, t')|_{t=t_j, t'=t_{j'}}$ .
- $\epsilon$ : Small jitter value for numerical stability ( $\epsilon \approx 10^{-6}$ ).
- $C_{inv,d} = (C_{\phi_d} + \epsilon I)^{-1}$ : Inverse of jittered prior covariance matrix.
- $m_{\phi_d} = C'_{\phi_d} C_{inv,d}$ : Matrix relating  $X_d$  to the conditional mean of its derivative  $\dot{X}_d$ .  $\mathbb{E}[\dot{X}_d | X_d] = m_{\phi_d} X_d$  (assuming zero prior mean).

- $K_{\phi_d} = C''_{\phi_d} - m_{\phi_d}(C'_{\phi_d})^\top$ : Conditional covariance matrix  $\mathbb{V}[\dot{X}_d|X_d]$ .
- $K_{inv,d} = (K_{\phi_d} + \epsilon I)^{-1}$ : Inverse of jittered conditional covariance matrix (precision matrix).
- $M^{Band}$ : Banded matrix approximation of matrix  $M$  with bandwidth  $b$ .
- $\beta = (\beta_{deriv}, \beta_{level}, \beta_{obs})$ : Prior temperature vector.
- $\Psi$ : Full parameter vector for sampling, e.g.,  $(\text{vec}(X), \theta, \log \sigma)$  or  $(\text{vec}(X), \theta)$ .
- $\|\mathbf{v}\|_{\mathbf{A}}^2 = \mathbf{v}^\top \mathbf{A} \mathbf{v}$ : Squared Mahalanobis norm.

## 1 Initialization

The goal of initialization is to determine suitable starting values for the latent states ( $X^{(0)}$ ), ODE parameters ( $\theta^{(0)}$ ), observation noise standard deviations ( $\sigma^{(0)}$ ), and GP hyperparameters ( $\phi^{(0)}$ ). These serve as the initial state  $\Psi^{(0)}$  for the MCMC sampler.

### 1.1 Handling User-Provided Initial State

If a complete initial state vector  $\Psi_{init}$  is provided by the user, it is unpacked directly into  $X^{(0)}$ ,  $\theta^{(0)}$ , and possibly the log-transformed  $\log \sigma^{(0)}$ . The parameter vector  $\theta^{(0)}$  is clamped to lie within the specified bounds  $\Theta = [\theta_{lower}, \theta_{upper}]$ . The flag `sigma_is_fixed` is determined based on whether  $\Psi_{init}$  includes the  $\log \sigma$  component. If  $\sigma$  is sampled (i.e., `sigma_is_fixed` is false), the initial value is set as  $\sigma^{(0)} = \exp(\log \sigma^{(0)})$ . If  $\sigma$  is fixed,  $\sigma^{(0)}$  is set to the value  $\sigma_{fixed}$  provided in the configuration (this is mandatory if  $\Psi_{init}$  implies fixed sigma). Similarly, if  $\Psi_{init}$  implies fixed sigma, the GP hyperparameters  $\phi_{fixed}$  must be provided in the configuration, and  $\phi^{(0)}$  is set to  $\phi_{fixed}$ . If  $\Psi_{init}$  is provided, the subsequent initialization steps for  $\phi, \sigma, X, \theta$  are skipped, and the algorithm proceeds directly to pre-computing the GP covariance structures using  $\phi^{(0)}$ .

### 1.2 Initialize GP Hyperparameters ( $\phi$ ) and Noise ( $\sigma$ )

This step is performed if  $\Psi_{init}$  was not provided, AND either the GP hyperparameters  $\phi_{fixed}$  were not provided or the noise standard deviations  $\sigma_{fixed}$  were not provided (meaning  $\sigma$  needs to be sampled).

The determination of whether  $\sigma$  is fixed relies on the configuration: `sigma_is_fixed` is true if and only if both  $\sigma_{fixed}$  and  $\phi_{fixed}$  are provided.

The initialization proceeds dimension by dimension ( $d = 1, \dots, D$ ):

1. Identify the valid (non-NaN) observations  $Y_{valid,d}$  in the  $d$ -th column of  $Y$  and their corresponding times  $t_{valid,d}$ .
2. Compute initial guesses for the log-transformed hyperparameters:  $(\log \phi_{d,1}^{(guess)}, \log \phi_{d,2}^{(guess)})$  and  $\log \sigma_d^{(guess)}$ . These guesses are derived from basic statistics of  $Y_{valid,d}$  (e.g., variance, median absolute deviation) and the time range of observations. If  $\phi_{fixed}$  was provided (but  $\sigma_{fixed}$  was not), its values are used for the  $\phi$  guess:  $(\log \phi_{fixed,1,d}, \log \phi_{fixed,2,d})$ .
3. Define the Negative Log Marginal Likelihood (NLML) objective function for dimension  $d$ . Assuming a GP prior  $x_d \sim \mathcal{GP}(0, \mathcal{K}_{\phi_d})$  and observation model  $y_d = x_d + \epsilon_d$  with  $\epsilon_d \sim \mathcal{N}(0, \sigma_d^2 I)$ , the marginal likelihood integrates out  $x_d$ . The NLML, as a function of log-parameters  $\lambda = (\log \phi_{d,1}, \log \phi_{d,2}, \log \sigma_d)$ , is:

$$L_{NLML}(\lambda_1, \lambda_2, \lambda_3) = \frac{1}{2} \log |\mathbf{K}_{\phi_d^*} + (\sigma_d^*)^2 \mathbf{I}| + \frac{1}{2} \mathbf{Y}_{valid,d}^\top (\mathbf{K}_{\phi_d^*} + (\sigma_d^*)^2 \mathbf{I})^{-1} \mathbf{Y}_{valid,d} + \text{const}$$

where  $\phi_{d,1}^* = \exp(\lambda_1)$ ,  $\phi_{d,2}^* = \exp(\lambda_2)$ ,  $\sigma_d^* = \exp(\lambda_3)$ , and  $\mathbf{K}_{\phi_d^*} = [\mathcal{K}_{(\phi_{d,1}^*, \phi_{d,2}^*)}(t, t')]_{t, t' \in t_{valid,d}}$ .

4. Minimize  $L_{NLML}(\lambda)$  using a numerical optimization algorithm (e.g., Nelder-Mead), starting from the initial guess  $\lambda^{(guess)}$ , to obtain the optimized log-parameters  $\lambda^* = (\log \phi_{d,1}^*, \log \phi_{d,2}^*, \log \sigma_d^*)$ .
5. Store the optimized parameters:
  - If  $\phi_{fixed}$  was not provided, store  $\phi_{est,1,d} = \exp(\lambda_1^*)$  and  $\phi_{est,2,d} = \exp(\lambda_2^*)$ .
  - If `sigma_is_fixed` is false, store  $\sigma_{est,d} = \exp(\lambda_3^*)$ .

After iterating through all dimensions, set the final initial GP hyperparameters  $\phi^{(0)}$  to  $\phi_{est}$  (if estimated) or  $\phi_{fixed}$  (if provided). Set the initial noise  $\sigma^{(0)}$  to  $\sigma_{est}$  (if estimated) or  $\sigma_{fixed}$  (if provided and fixed).

### 1.3 Initialize Latent States ( $X$ )

This step is performed if  $\Psi_{init}$  and  $X_{init}$  were not provided. Initialize the latent state matrix  $X^{(0)} \in \mathbb{R}^{n \times D}$ . For each dimension  $d = 1, \dots, D$ , identify the non-NaN observations  $Y_{valid,d}$  at times  $t_{valid,d}$ . Linearly interpolate these values onto the full discretization grid  $\mathcal{I}$  to obtain the initial trajectory  $X_{:,d}^{(0)}$ . Handle cases with fewer than two observations appropriately (e.g., constant extrapolation).

If  $X_{init}$  was provided, set  $X^{(0)} = X_{init}$  after checking dimensions.

### 1.4 Initialize ODE Parameters ( $\theta$ )

This step is performed if  $\Psi_{init}$  and  $\theta_{init}$  were not provided. Initialize the ODE parameter vector  $\theta^{(0)} \in \mathbb{R}^k$ . Typically, this is done by taking the midpoint of the provided bounds  $\Theta = [\theta_{lower}, \theta_{upper}]$ , or offsetting slightly from a bound if one side is infinite.

If  $\theta_{init}$  was provided, set  $\theta^{(0)} = \theta_{init}$ .

Finally, clamp  $\theta^{(0)}$  to ensure it lies within the bounds:  $\theta^{(0)} = \max(\theta_{lower}, \min(\theta^{(0)}, \theta_{upper}))$ .

### 1.5 Pre-compute GP Covariance Structures

Using the initialized or provided GP hyperparameters  $\phi^{(0)}$ , pre-compute the necessary covariance matrices and their banded approximations for each dimension  $d = 1, \dots, D$ . These matrices are essential for evaluating the GP prior terms ( $\log p(X_d)$ ) and the GP conditional derivative terms ( $\log p(\dot{X}_d = F_d | X_d)$ ) in the log-likelihood. Store these in a structure, `gp_cov_all_dims`.

For each dimension  $d$ :

1. Select the kernel function  $\mathcal{K}_{\phi_d}$  based on the kernel type specified in the configuration and the parameters  $\phi_d = \phi_{:,d}^{(0)}$ .
2. Compute the dense  $n \times n$  prior covariance matrix  $C_{\phi_d} = [\mathcal{K}_{\phi_d}(t_j, t_{j'})]_{j,j'=1}^n$ .
3. Compute the dense first time derivative matrix  $C'_{\phi_d}$ .
4. Compute the dense second time derivative matrix  $C''_{\phi_d}$ .
5. Compute the inverse of the jittered prior covariance:  $C_{inv,d} = (C_{\phi_d} + \epsilon I)^{-1}$ . This is needed for the  $\mathcal{L}_{level}$  term and its gradient.
6. Compute the matrix  $m_{\phi_d} = C'_{\phi_d} C_{inv,d}$ . This relates the level  $X_d$  to the conditional mean of the derivative  $\mathbb{E}[\dot{X}_d | X_d] = m_{\phi_d} X_d$  (assuming zero prior mean).
7. Compute the conditional covariance matrix  $K_{\phi_d} = \mathbb{V}[\dot{X}_d | X_d] = C''_{\phi_d} - m_{\phi_d} (C'_{\phi_d})^\top$ .
8. Compute the inverse of the jittered conditional covariance (precision matrix):  $K_{inv,d} = (K_{\phi_d} + \epsilon I)^{-1}$ . This is needed for the  $\mathcal{L}_{deriv}$  term and its gradient.

9. Create banded matrix approximations using the specified bandwidth  $b$ :  $C_{inv,d}^{Band}$ ,  $m_{\phi_d}^{Band}$ ,  $K_{inv,d}^{Band}$ . These are used for efficient computation.
10. Store  $C_{inv,d}^{Band}$ ,  $m_{\phi_d}^{Band}$ ,  $K_{inv,d}^{Band}$  in the structure `gp_cov_all_dims[d]`. (The dense matrices  $K_{\phi_d}$ ,  $C_{\phi_d}$  etc. may also be stored for reference).

## 1.6 Construct Initial Sampler State $\Psi^{(0)}$

This step constructs the final initial state vector for the MCMC sampler, unless it was provided directly as  $\Psi_{init}$ .

- If `sigma_is_fixed` is true:

$$\Psi^{(0)} = \begin{pmatrix} \text{vec}(X^{(0)}) \\ \theta^{(0)} \end{pmatrix}$$

- If `sigma_is_fixed` is false: Compute the log-transformed noise  $\log \sigma^{(0)} = \log(\max(\sigma^{(0)}, \text{small\_val}))$  (element-wise, ensuring positivity).

$$\Psi^{(0)} = \begin{pmatrix} \text{vec}(X^{(0)}) \\ \theta^{(0)} \\ \log \sigma^{(0)} \end{pmatrix}$$

This  $\Psi^{(0)}$  is the starting point for the MCMC sampler described in Section 3.

## 2 Log-Likelihood & Gradient Calculation: $\mathcal{L}(\Psi)$ and $\nabla \mathcal{L}(\Psi)$

This section describes the calculation of the log-posterior density  $\mathcal{L}(\Psi)$  and its gradient  $\nabla \mathcal{L}(\Psi)$  with respect to the full parameter vector  $\Psi = (\text{vec}(X), \theta, \dots)$ . The log-posterior combines the log-likelihood of the observations, the log-prior density from the GP assumptions on levels  $X$  and derivatives  $\dot{X}$ , and the log-prior density of the ODE parameters  $\theta$  (often assumed uniform within bounds, contributing only via the bounds). The gradient is required for HMC sampling.

### 2.1 Parameter Unpacking and Transformation

Given an input parameter vector  $\Psi$ :

1. Unpack  $\Psi$  into the latent state matrix  $X \in \mathbb{R}^{n \times D}$  and the ODE parameter vector  $\theta \in \mathbb{R}^k$ .
2. Check if  $\sigma$  is fixed using the `sigma_is_fixed` flag determined during initialization.
3. If `sigma_is_fixed` is true, use the fixed value  $\sigma = \sigma^{(0)}$  determined during initialization. The log-Jacobian contribution is  $\mathcal{L}_{jac} = 0$ .
4. If `sigma_is_fixed` is false, unpack the remaining elements of  $\Psi$  as  $\log \sigma \in \mathbb{R}^D$ . Transform to the original scale:  $\sigma = \exp(\log \sigma)$  (element-wise). Ensure resulting  $\sigma_d > 0$ . The log-Jacobian determinant for this transformation is  $\mathcal{L}_{jac} = \sum_{d=1}^D (\log \sigma)_d$ . This term is added to the total log-likelihood, effectively corresponding to a  $p(\sigma_d) \propto 1/\sigma_d$  prior on  $\sigma_d$ .

### 2.2 Log-Likelihood Term Calculation

Initialize the total log-likelihood  $\mathcal{L} = \mathcal{L}_{jac}$ . Pre-compute the ODE function values  $F \in \mathbb{R}^{n \times D}$  where  $F_{j,d} = f_d(X_{j,:}, \theta, t_j)$  for all  $j = 1..n, d = 1..D$ .

Iterate through each dimension  $d = 1, \dots, D$ :

1. Retrieve the pre-computed banded matrices  $C_{inv,d}^{Band}$ ,  $m_d^{Band}$ ,  $K_{inv,d}^{Band}$  from `gp_cov_all_dims[d]`.

2. Define vectors for the current dimension:  $X_d = X_{:,d}$ ,  $F_d = F_{:,d}$ ,  $Y_d = Y_{:,d}$ .
3. Identify the set of time indices  $\mathcal{V}_d = \{j | Y_{j,d} \text{ is not NaN}\}$  where observations are available for dimension  $d$ .
4. Calculate the level fit error:  $\text{fitLevelError}_d = X_d - Y_d$ . Set entries corresponding to NaN observations in  $Y_d$  to zero.
5. Calculate the expected derivative based on the GP mean:  $\text{mphi\_x}_d = m_d^{\text{Band}} X_d$ .
6. Calculate the derivative fit error:  $\text{fitDerivError}_d = F_d - \text{mphi\_x}_d$ .
7. Calculate intermediate terms needed for the likelihood components:
  - $\text{Kinv\_fitDerivError}_d = K_{\text{inv},d}^{\text{Band}} \times \text{fitDerivError}_d$
  - $\text{Cinv\_x}_d = C_{\text{inv},d}^{\text{Band}} \times X_d$
8. Calculate the three components of the log-likelihood for dimension  $d$ :
  - **Level Log-Likelihood** ( $\log p(X_d)$ ): The GP prior assumes  $X_d \sim \mathcal{N}(0, C_{\phi_d})$ . The log-density is:

$$\log p(X_d) = -\frac{1}{2} X_d^\top C_{\phi_d}^{-1} X_d - \frac{1}{2} \log |C_{\phi_d}| - \frac{n}{2} \log(2\pi)$$

Using the pre-computed inverse of the jittered matrix  $C_{\text{inv},d}$  and the term  $\text{Cinv\_x}_d = C_{\text{inv},d} X_d$ , and ignoring constant terms and log-determinants (if  $\phi$  is fixed), this component is calculated as:

$$\mathcal{L}_{\text{level},d} = -\frac{1}{2} X_d^\top C_{\text{inv},d} X_d = -\frac{1}{2} X_d^\top (\text{Cinv\_x}_d)$$

This penalizes deviations from the zero mean prior, weighted by the prior precision  $C_{\text{inv},d}$ , corresponding to  $-\frac{1}{2} \mathcal{L}_{GP,d}$  in Eq. (5) of the reference PDF [71-73].

- **Observation Log-Likelihood** ( $\log p(Y_d | X_d, \sigma_d^2)$ ): The model assumes  $Y_{j,d} | X_{j,d} \sim \mathcal{N}(X_{j,d}, \sigma_d^2)$  independently for observed points  $j \in \mathcal{V}_d$ . The joint log-density is:

$$\begin{aligned} \log p(Y_d(\mathcal{V}_d) | X_d(\mathcal{V}_d), \sigma_d^2) &= \sum_{j \in \mathcal{V}_d} \log p(Y_{j,d} | X_{j,d}, \sigma_d^2) \\ &= \sum_{j \in \mathcal{V}_d} \left[ -\frac{1}{2} \log(2\pi\sigma_d^2) - \frac{(X_{j,d} - Y_{j,d})^2}{2\sigma_d^2} \right] \\ &= -\frac{|\mathcal{V}_d|}{2} \log(2\pi\sigma_d^2) - \frac{1}{2\sigma_d^2} \sum_{j \in \mathcal{V}_d} (X_{j,d} - Y_{j,d})^2 \end{aligned}$$

Using the pre-calculated  $\text{fitLevelError}_{d,j} = X_{j,d} - Y_{j,d}$ , this is computed as:

$$\mathcal{L}_{\text{obs},d} = -\frac{1}{2\sigma_d^2} \sum_{j \in \mathcal{V}_d} (\text{fitLevelError}_{d,j})^2 - \frac{|\mathcal{V}_d|}{2} \log(2\pi\sigma_d^2)$$

This represents the goodness-of-fit, corresponding to  $-\frac{1}{2} \mathcal{L}_{\text{Obs},d}$  in Eq. (5) [73].

- **Derivative Log-Likelihood** ( $\log p(\dot{X}_d = F_d | X_d)$ ): From GP properties, the conditional distribution is  $\dot{X}_d | X_d \sim \mathcal{N}(m_{\phi_d} X_d, K_{\phi_d})$ . The manifold constraint implies  $\dot{X}_d = F_d$ . Evaluating the log-density of this conditional normal distribution at  $F_d$ :

$$\log p(F_d | X_d) = -\frac{1}{2} (F_d - m_{\phi_d} X_d)^\top K_{\phi_d}^{-1} (F_d - m_{\phi_d} X_d) - \frac{1}{2} \log |K_{\phi_d}| - \frac{n}{2} \log(2\pi)$$

Using the jittered inverse  $K_{inv,d} = (K_{\phi_d} + \epsilon I)^{-1}$  and the derivative error  $E_d = \text{fitDerivError}_d = F_d - \text{mphi}_d X_d$ , and ignoring constants and log-determinants, this component is calculated as:

$$\mathcal{L}_{deriv,d} = -\frac{1}{2} E_d^\top K_{inv,d} E_d = -\frac{1}{2} (\text{fitDerivError}_d)^\top (\text{Kinv\_fitDerivError}_d)$$

This term enforces the ODE manifold constraint by penalizing the squared Mahalanobis distance between the ODE dynamics  $F_d$  and the GP's expected derivative  $m_{\phi_d} X_d$ , weighted by the conditional precision  $K_{inv,d}$ . It corresponds to  $-\frac{1}{2} \mathcal{L}_{Constraint,d}$  in Eq. (5) [74-77].

9. Add the weighted components to the total log-likelihood:

$$\mathcal{L} \leftarrow \mathcal{L} + \frac{\mathcal{L}_{obs,d}}{\beta_{obs}} + \frac{\mathcal{L}_{deriv,d}}{\beta_{deriv}} + \frac{\mathcal{L}_{level,d}}{\beta_{level}}$$

### 2.3 Gradient Calculation

The gradient  $\nabla \mathcal{L}(\Psi)$  required for HMC is computed by summing the gradients of the weighted log-likelihood components derived above. These gradients are obtained analytically using standard matrix and vector calculus.

**Analytical Gradient Components** The gradient of the total log-likelihood is the sum of the gradients of its constituent parts (weighted by inverse temperatures  $1/\beta$ ), plus the gradient of the log-Jacobian term  $\mathcal{L}_{jac}$  (if applicable). We consider the gradient of each main term:

- **Gradient of  $\mathcal{L}_{level,d}$ :** Since  $\mathcal{L}_{level,d} = -\frac{1}{2} X_d^\top C_{inv,d} X_d + \text{const}$ , its gradient w.r.t.  $X_d$  is  $\nabla_{X_d} \mathcal{L}_{level,d} = -C_{inv,d} X_d$ . It has no dependence on  $\theta$  or  $\sigma$ .
- **Gradient of  $\mathcal{L}_{obs,d}$ :** Since  $\mathcal{L}_{obs,d} = -\frac{1}{2\sigma_d^2} \sum_{j \in \mathcal{V}_d} (X_{j,d} - Y_{j,d})^2 + \dots$ , its gradient w.r.t.  $X_{j,d}$  is  $\nabla_{X_{j,d}} \mathcal{L}_{obs,d} = -\frac{1}{\sigma_d^2} (X_{j,d} - Y_{j,d})$  for observed points ( $j \in \mathcal{V}_d$ ) and zero otherwise. The gradient w.r.t.  $\sigma_d$  is  $\nabla_{\sigma_d} \mathcal{L}_{obs,d} = \frac{1}{\sigma_d^3} \sum_{j \in \mathcal{V}_d} (X_{j,d} - Y_{j,d})^2 - \frac{|\mathcal{V}_d|}{\sigma_d}$ . It does not depend on  $\theta$ .
- **Gradient of  $\mathcal{L}_{deriv,d}$ :** This term,  $\mathcal{L}_{deriv,d} = -\frac{1}{2} (F_d - m_{\phi_d} X_d)^\top K_{inv,d} (F_d - m_{\phi_d} X_d) + \text{const}$ , depends on  $X$  (directly via  $m_{\phi_d} X_d$  and indirectly via  $F_d = f(X, \theta, t)$ ) and  $\theta$  (via  $F_d$ ). Let  $E_d = F_d - m_{\phi_d} X_d$  be the error vector. Using the chain rule for quadratic forms ( $\nabla_z (-\frac{1}{2} v^\top A v) = -(\nabla_z v)^\top A v$  for symmetric  $A$ ), we get:
  - $\nabla_\theta \mathcal{L}_{deriv,d} = -(\nabla_\theta E_d)^\top K_{inv,d} E_d$ . Since only  $F_d$  depends on  $\theta$ ,  $\nabla_\theta E_d = \nabla_\theta F_d$ , where the  $(j, p)$ -th element of  $\nabla_\theta F_d$  is  $\frac{\partial f_d(X_{j,:}, \theta, t_j)}{\partial \theta_p}$ . Thus,  $\nabla_\theta \mathcal{L}_{deriv,d} = -(\nabla_\theta F_d)^\top K_{inv,d} (F_d - m_{\phi_d} X_d)$ .
  - $\nabla_X \mathcal{L}_{deriv,d} = -(\nabla_X E_d)^\top K_{inv,d} E_d$ . Here,  $\nabla_X E_d = \nabla_X F_d - \nabla_X (m_{\phi_d} X_d)$ .  $\nabla_X F_d$  involves the ODE state Jacobian  $J_x$ , specifically  $\frac{\partial F_{j,d}}{\partial X_{j',d'}} = \delta_{jj'} J_x(j)[d, d']$ . The gradient  $\nabla_X (m_{\phi_d} X_d)$  involves  $m_{\phi_d}^\top$ . Combining these leads to contributions to  $\nabla_X \mathcal{L}_{j',d'}$  involving both  $J_x(j')[d, d']$  and  $(m_{\phi_d})^\top$ .

This term does not depend directly on  $\sigma$ .

- **Gradient of  $\mathcal{L}_{jac}$ :** If  $\sigma$  is sampled,  $\mathcal{L}_{jac} = \sum \log \sigma_d$ . Then  $\nabla_{\log \sigma} \mathcal{L}_{jac} = \mathbf{1}$ , and gradients w.r.t  $X, \theta$  are zero.

The subsequent steps detail how these analytical gradient components are computed and accumulated efficiently using the pre-calculated matrices and Jacobians.

Initialize gradient components:  $\nabla_X \mathcal{L} = \mathbf{0} \in \mathbb{R}^{n \times D}$ ,  $\nabla_\theta \mathcal{L} = \mathbf{0} \in \mathbb{R}^k$ ,  $\nabla_\sigma \mathcal{L} = \mathbf{0} \in \mathbb{R}^D$ .

#### Accumulating Gradients:

1. **Loop through dimensions** ( $d = 1..D$ ): Calculate contributions independent of ODE coupling.

- $\nabla_X \mathcal{L}_{:,d} \leftarrow \nabla_X \mathcal{L}_{:,d} + \frac{1}{\beta_{level}} (\nabla_{X_d} \mathcal{L}_{level,d}) = \nabla_X \mathcal{L}_{:,d} - \frac{1}{\beta_{level}} C_{inv,d} X_d$
- For  $j \in \mathcal{V}_d$ :  $\nabla_X \mathcal{L}_{j,d} \leftarrow \nabla_X \mathcal{L}_{j,d} + \frac{1}{\beta_{obs}} (\nabla_{X_{j,d}} \mathcal{L}_{obs,d}) = \nabla_X \mathcal{L}_{j,d} - \frac{1}{\beta_{obs} \sigma_d^2} (\text{fitLevelError}_{d,j})$
- $\nabla_X \mathcal{L}_{:,d} \leftarrow \nabla_X \mathcal{L}_{:,d} + \frac{1}{\beta_{deriv}} (\nabla_{X_d} \mathcal{L}_{deriv,d})_{\text{mphi}} = \nabla_X \mathcal{L}_{:,d} + \frac{1}{\beta_{deriv}} (m_d^{Band})^\top (\text{Kinv\_fitDerivError}_d)$   
(From  $-m_{\phi_d} X_d$  part of  $\mathcal{L}_{deriv,d}$ )
- If not **sigma\_is\_fixed**:  $\nabla_\sigma \mathcal{L}_d \leftarrow \nabla_\sigma \mathcal{L}_d + \frac{1}{\beta_{obs}} (\nabla_{\sigma_d} \mathcal{L}_{obs,d}) = \nabla_\sigma \mathcal{L}_d + \frac{1}{\beta_{obs}} \left( \frac{\text{SSE}_d}{\sigma_d^3} - \frac{|\mathcal{V}_d|}{\sigma_d} \right)$

2. **Loop through time** ( $j = 1..n$ ): Calculate contributions involving ODE coupling from the  $F_d$  term in  $\mathcal{L}_{deriv}$ .

- Compute Jacobians  $J_x(j) = \nabla_x f(X_{j,:}, \theta, t_j)$  and  $J_\theta(j) = \nabla_\theta f(X_{j,:}, \theta, t_j)$ .
- For each dimension  $d = 1..D$  (of the derivative error term):
  - Let  $KFE_{d,j} = [\text{Kinv\_fitDerivError}_d]_j$ .
  - Accumulate gradient w.r.t.  $X_{j,d'}$  ( $d' = 1..D$ ):

$$\nabla_X \mathcal{L}_{j,d'} \leftarrow \nabla_X \mathcal{L}_{j,d'} - \frac{1}{\beta_{deriv}} J_x(j)[d, d'] \cdot KFE_{d,j}$$

(This corresponds to the  $(\nabla_X F_d)^\top (-K_{inv,d} E_d)$  part, summed over  $d$ ).

- Accumulate gradient w.r.t.  $\theta$ :

$$\nabla_\theta \mathcal{L} \leftarrow \nabla_\theta \mathcal{L} - \frac{1}{\beta_{deriv}} J_\theta(j)[d, :]^\top \cdot KFE_{d,j}$$

(This corresponds to the  $(\nabla_\theta F_d)^\top (-K_{inv,d} E_d)$  part, summed over  $d$ ).

#### Assemble Final Gradient Vector $\nabla \mathcal{L}(\Psi)$ :

1. Combine the accumulated  $\nabla_X \mathcal{L}$  (flattened) and  $\nabla_\theta \mathcal{L}$  into the appropriate parts of the final gradient vector  $\nabla \mathcal{L}$ .
2. If **sigma\_is\_fixed** is false, calculate the gradient with respect to the sampled variable  $\log \sigma$ :

$$\nabla_{\log \sigma} \mathcal{L} = (\nabla_\sigma \mathcal{L} \odot \sigma) + \mathbf{1}$$

(The first term uses the chain rule, the second is the gradient of  $\mathcal{L}_{jac}$ ). Assign this  $\nabla_{\log \sigma} \mathcal{L}$  to the final elements of  $\nabla \mathcal{L}$ .

The function returns the pair  $(\mathcal{L}, \nabla \mathcal{L})$ .

## 3 Sampling

The final stage employs a Markov Chain Monte Carlo (MCMC) algorithm, typically the No-U-Turn Sampler (NUTS), a variant of Hamiltonian Monte Carlo (HMC), to draw samples from the posterior distribution  $p(\Psi|Y) \propto \exp(\mathcal{L}(\Psi))$ .

### 3.1 Sampler Setup

1. Initialize the NUTS sampler, providing the log-likelihood function  $\mathcal{L}(\Psi)$  and its gradient  $\nabla\mathcal{L}(\Psi)$  calculated as described in Section 2.
2. Define the HMC metric  $M$ , which represents the mass matrix (often diagonal Euclidean, meaning  $M$  is diagonal). This influences the kinetic energy term in the Hamiltonian.
3. Define the numerical integrator (typically Leapfrog) used to simulate Hamiltonian dynamics, along with an initial step size  $\epsilon_0$ .
4. Configure an adaptor (e.g., StanHMCAdaptor) responsible for tuning the step size  $\epsilon$  and potentially the metric  $M$  during the initial "burn-in" or "warmup" phase of the MCMC run. The goal is often to achieve a target acceptance ratio (e.g., 0.8).

### 3.2 MCMC Iteration Loop

Set the current state of the chain  $\Psi_{current} = \Psi^{(0)}$  (from Section 1). Initialize storage  $\mathcal{S}$  for accepted samples. Iterate  $i$  from 1 to  $N_{iter}$ :

1. Sample momentum vector  $p$  from its conditional distribution, typically  $p \sim \mathcal{N}(0, M)$ .
2. Propose a new state  $(\Psi^*, p^*)$  by simulating the Hamiltonian dynamics  $H(\Psi, p) = -\mathcal{L}(\Psi) + \frac{1}{2}p^\top M^{-1}p$  starting from  $(\Psi_{current}, p)$ . The NUTS algorithm dynamically determines the simulation length (number of Leapfrog steps) to avoid random walks and efficiently explore the parameter space.
3. Calculate the Metropolis-Hastings acceptance probability:

$$\alpha = \min(1, \exp(H(\Psi_{current}, p) - H(\Psi^*, p^*)))$$

$$\alpha = \min\left(1, \exp\left(\mathcal{L}(\Psi^*) - \mathcal{L}(\Psi_{current}) + \frac{1}{2}p^\top M^{-1}p - \frac{1}{2}(p^*)^\top M^{-1}p^*\right)\right)$$

4. Sample a random number  $u \sim \text{Uniform}(0, 1)$ .
5. If  $u < \alpha$ , accept the proposal:  $\Psi_{current} = \Psi^*$ . Otherwise, reject:  $\Psi_{current}$  remains unchanged.
6. If the current iteration  $i$  is after the burn-in phase ( $i > N_{burn}$ ), store the current state  $\Psi_{current}$  in the sample collection  $\mathcal{S}$ .
7. If the current iteration  $i$  is within the burn-in phase ( $i \leq N_{burn}$ ), use the outcome of the step (acceptance statistic, proposed state) to update the step size  $\epsilon$  and potentially the metric  $M$  via the configured adaptor.

### 3.3 Post-processing

After the MCMC loop completes, process the collected samples  $\Psi^{(i)} \in \mathcal{S}$  (for  $i = N_{burn} + 1, \dots, N_{iter}$ ):

1. For each sample  $\Psi^{(i)}$ , unpack it into its components:  $X^{(i)} \in \mathbb{R}^{n \times D}$ ,  $\theta^{(i)} \in \mathbb{R}^k$ .
2. Determine  $\sigma^{(i)}$ :
  - If `sigma_is_fixed` was true during the run, set  $\sigma^{(i)} = \sigma^{(0)}$  (the fixed value).
  - If `sigma_is_fixed` was false, extract the log  $\sigma^{(i)}$  component from  $\Psi^{(i)}$  and transform it back:  $\sigma^{(i)} = \exp(\log \sigma^{(i)})$ .



The final output of the algorithm is the set of post-burn-in samples  $\{(X^{(i)}, \theta^{(i)}, \sigma^{(i)})\}_{i=N_{burn}+1}^{N_{iter}}$ . These samples represent draws from the posterior distribution and can be used for inference (e.g., calculating means, credible intervals).